


Governing Hidden Inference: Output-Centric Privacy Control for AI Apps

Sieun Park, Youngki Lee | Seoul National University



Problem: Input Permissions Leave an Inference Gap




Blocking access does not stop inference
: denied attributes can reappear in AI-generated outputs.
→ Input permissions cannot govern inferred outputs!




Why are those flags hanging there?
This appears to be **Green Street, Cambridge, UK**.
They are decorative bunting used to create a festive atmosphere and attract visitors.
(How did it know where this was..?
The photo didn't have any signs..)

Key Idea: Output-Centric Privacy Policy over Derived Attributes


Goal: **Align** AI app outputs with **users' disclosure intent**
: mediate derived attributes before externalization according to user policy



Preliminary Study: Privacy Risk Shifts Across the AI Pipeline



Ongoing Work



Contact

Sieun Park
Watch the project video or
leave questions and
feedback via the QR code!

